

Lab 6



TA: 何國豪、鄭余玄

11/15 12:00

Lab6 Deadline


no demo

Outline

1. Solve **CartPole-v1** using **DQN**
2. Solve **Pendulum-v0** using **DDPG**
3. Modify and Run Sample Code
4. Scoring Criteria
5. Reminders
6. Sample Code Guides

1. CartPole-v1

- Observation [4]
 - Cart Position
 - Cart Velocity
 - Pole Angle
 - Pole Velocity at Tip
- Action [2]
 - Left
 - Right
- Reward
 - **+1** for every time step



- Starting State
 - All observations are assigned a uniform random value between ± 0.05
- Episode Termination
 - Pole Angle is more than $\pm 12^\circ$
 - Center of the cart reaches the edge of the display
 - Episode length is greater than **500**

Deep Q-Network (DQN)

Algorithm 1: deep Q-learning with experience replay.

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

With probability ϵ select a random action a_t
otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

Execute action a_t in emulator and observe reward r_t and image x_{t+1}

Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Every C steps reset $\hat{Q} = Q$

End For


End For

TODO:

- Construct the neural network
- Select action according to epsilon-greedy
- Construct Q-values and target Q-values
- Calculate loss function
- Update behavior and target network
- Understand deep Q-learning mechanisms

2. Pendulum-v0

- Observation [3]
 - $\cos(\theta)$ (Angle)
 - $\sin(\theta)$ (Angle)
 - $\theta_{\dot{}}$ (Angular Velocity)
- Action [1]
 - Joint Effort
- Reward
 - In essence, the goal is to remain at **zero angle** (vertical), with the least rotational velocity, and the least effort.
- Starting State
 - Random angle from $-\pi$ to π , and random velocity between -1 and 1
- Episode Termination
 - Episode length is greater than 200



Deep Deterministic Policy Gradient (DDPG)

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q$, $\theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M **do**

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation state s_1

for $t = 1, T$ **do**

 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in R

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'})|\theta^{Q'})$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

 Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

TODO:

- Construct neural networks of both actor and critic
- Select action according to the actor and the exploration noise
- Update critic
- Update actor
- Update target network softly
- Understand the mechanism of actor-critic

end for
end for

3. Modify Sample Code

1. Find a `#TODO` comment with hints
2. remove the `raise NotImplementedError`

3. Run Sample Code

- Simply train and test: `python dqn.py`
- Only test and render: `python dqn.py --restore --render`
- Help message: `python dqn.py --help`
- Other usages:
 - Save as different model name: `python dqn.py -m cart_model`
 - Load and test different model: `python dqn.py -m cart_model --restore`
 - Only run on cpu: `python dqn.py -d cpu`
 - Only run on 3rd gpu: `python dqn.py -d cuda:2`
 - Set episode length to 2000: `python dqn.py -e 2000`

4. Scoring Criteria

Show your work, otherwise no credit will be granted.

- Report (80%)
 - (DO **explain**; do not copy and paste your codes.)
- Report Bonus (10%)
 - Implement and perform experiments on **Double-DQN**. (5%)
 - **Extra experiments**; e.g., random process comparison (5%)
- Performance (20%)
 - [**CartPole-v1**] Average reward of 10 testing episodes: $\text{Average} \div 5$
 - [Pendulum-v0] Average reward of 10 testing episodes: $(\text{Average} + 700) \div 5$

5. Reminders

- Your network architecture and hyper-parameters **can** differ from the defaults.
- Be careful of the **target q-value** at the end of an episode in DQN.
- Ensure the **shape** of tensors all the time especially when calculating the **loss**.
- `with no_grad():` scope is the same as `xxx.detach()`
- Be aware of the **indentation** of hints.
- When testing DDPG, action selection need **NOT** include the noise.

Sample Code

Sample Code

- `dqn-example.py`
- `ddpg-example.py`

Replay Memory

```
class ReplayMemory:
    def __init__(self, capacity):
        self._buffer = deque(maxlen=capacity)
```

```
    def __len__(self):
        return len(self._buffer)
```

```
    def append(self, *transition):
        # (state, action, reward, next_state, done)
        self._buffer.append(tuple(map(tuple, transition)))
```

```
    def sample(self, batch_size=1):
        return random.sample(self._buffer, batch_size)
```

Usage in algorithm:

Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D
Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Greedy Action Selection

```
def select_action(epsilon, state, action_dim=2):  
    """epsilon-greedy based on behavior network"""  
    ## TODO ##  
    raise NotImplementedError
```

With probability ε select a random action a_t
otherwise select $a_t = \operatorname{argmax}_a Q(\phi(s_t), a; \theta)$

DQN

```
class DQN(nn.Module):  
    def __init__(self, state_dim=4, action_dim=2, hidden_dim=24):  
        super().__init__()  
        ## TODO ##  
        raise NotImplementedError  
  
    def forward(self, x):  
        ## TODO ##  
        raise NotImplementedError
```


Update Behavior Network

```
def update_behavior_network():  
    def transitions_to_tensors(transitions, device=args.device):  
        """convert a batch of transitions to tensors"""  
        return (torch.Tensor(x).to(device) for x in zip(*transitions))
```

```
# sample a minibatch of transitions
```

```
transitions = memory.sample(args.batch_size)
```

```
state, action, reward, next_state, done = transitions_to_tensors(transitions)
```

```
# TODO: loss
```

```
# q_value = ?
```

```
# with torch.no_grad():
```

```
#     q_next = ?
```

```
#     q_target = ?
```

```
# loss = criterion(q_value, q_target)
```

```
raise NotImplementedError
```

Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

Actor Action Selection

```
def select_action(state, low=-2, high=2):  
    """based on the behavior (actor) network and exploration noise"""  
    ## TODO ##  
    # with torch.no_grad():  
    #     action = ? + ?  
    #     return max(min(action, high), low)  
    raise NotImplementedError
```

Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

Random Process

```
class OrnsteinUhlenbeckProcess:
    """1-dimension Ornstein-Uhlenbeck process"""
    def sample(self, mu=0, std=.2, theta=.15, dt=1e-2, sqrt_dt=1e-1):
        self.x += theta * (mu - self.x) * dt + std * sqrt_dt * random.gauss(0, 1)
        return self.x

    def reset(self, x0=0):
        self.x = x0
```

used only during training

Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

References

1. Mnih, Volodymyr et al. “Playing Atari with Deep Reinforcement Learning.” ArXiv abs/1312.5602 (2013).
2. Mnih, Volodymyr et al. “Human-level control through deep reinforcement learning.” Nature 518 (2015): 529-533.
3. Van Hasselt, Hado, Arthur Guez, and David Silver. “Deep Reinforcement Learning with Double Q-Learning.” AAAI. 2016.
4. Lillicrap, Timothy P. et al. “Continuous control with deep reinforcement learning.” CoRR abs/1509.02971 (2015).
5. Silver, David et al. “Deterministic Policy Gradient Algorithms.” ICML (2014).
6. OpenAI. “OpenAI Gym Documentation.” Retrieved from Getting Started with Gym: <https://gym.openai.com/docs/>.
7. OpenAI. “OpenAI Wiki for Pendulum v0.” Retrieved from Github: <https://github.com/openai/gym/wiki/Pendulum-v0>.
8. PyTorch. “Reinforcement Learning (DQN) Tutorial.” Retrieved from PyTorch Tutorials: https://pytorch.org/tutorials/intermediate/reinforcement_q_learning.html.